

Aaryan Kale

+91 9820712368

aryankale1410@gmail.com

linkedin.com/in/aaryan-kale

github.com/aryankale14

Portfolio Website

Professional Summary

AI Engineer with experience building production-ready GenAI applications using LLMs, RAG, multi-agent systems, FastAPI, and PostgreSQL. Experienced in developing end-to-end AI products including research assistants, code intelligence platforms, and resume analysis tools. Passionate about scalable AI systems, backend engineering, and applied machine learning.

Skills

- **AI/ML:** Gemini, OpenAI, Amazon Bedrock, RAG, LangChain, LangGraph, Vector Search, Prompt Engineering, QLoRA
- **Databases:** PostgreSQL, pgvector, Pinecone, FAISS, SQLite
- **Languages & Web:** Python, SQL, Java, React
- **Backend:** FastAPI, Flask, REST APIs
- **Tools:** AWS (S3, EC2, Textract), Firebase, Supabase, Git, Render, Vercel, Netlify, Railway

Work Experience

Capgemini

Mumbai, October 2025 – Present

Analyst (GenAI & Intelligent Automation Initiative)

- Developed an Amazon Bedrock-powered AI agent for automated classification and remediation of Apache Airflow workflow failures, reducing manual incident triage time by 40% through LLM-driven recommendations.
- Worked with cross-functional teams to validate AI-generated remediation decisions and establish guardrails for production-safe automation.
- Optimized prompts and evaluation workflows for LLM-based incident classification, improving response quality and operational reliability.

Projects

PaperMind – Multi-Agent Deep Research Assistant

Feb 2026 – Mar 2026

Python, React, FastAPI, PostgreSQL (pgvector), Gemini AI, Firebase

Live: [PaperMind](#)

- Architected a multi-tenant RAG pipeline using a 13-step autonomous agent system to handle complex query planning, multi-hop reasoning, and document synthesis.
- Engineered scalable cloud vector storage using PostgreSQL and pgvector, successfully migrating from local FAISS indices to ensure persistent, isolated user data.
- Implemented a self-reflection loop with an evaluation agent to verify citation grounded-ness, effectively halting hallucinations before answer generation.
- Optimized hybrid retrieval (ANN dense vector + BM25 lexical search) to run efficiently in memory-constrained cloud environments without relying on heavy cross-encoders.

CodeRecall - Multi-Agent Map-Reduce Codebase Intelligence & RAG Platform

April 2026- May 2026

Python, React, FastAPI, PostgreSQL (pgvector), Gemini AI, Firebase

Live: [CodeRecall](#)

- Architected a distributed multi-agent Map-Reduce pipeline using FastAPI to analyze and summarize complex GitHub repositories concurrently, mitigating LLM context window degradation and eliminating output token limitations.
- Engineered a production-grade RAG infrastructure leveraging PostgreSQL (pgvector) and Gemini embeddings, utilizing advanced metadata filtering to deliver low-latency semantic code search and context-aware interactive Q&A.
- Built an automated multi-agent security auditing engine that statically maps high, medium, and low-priority vulnerabilities across complex codebases, compiling results into an interactive, globally synthesized security catalog.

HireReady – AI Resume Analyzer and Career Toolkit

Oct 2025 – Nov 2025

React, Flask, MongoDB, MySQL, Gemini API

Live: [HireReady](#)

- Built an AI-powered resume analysis platform with ATS scoring and job-description matching.
- Generated structured feedback including keyword gaps, strengths/weaknesses, and improvements using Gemini APIs.
- Generated ATS scores and keyword recommendations by comparing resumes against job descriptions.

Education

B.E. in Electronics and Telecommunication Engineering

Vidyalankar Institute of Technology, Mumbai

Nov 2022 – May 2025

Diploma in Electronics and Telecommunication Engineering

Bharati Vidyapeeth's Institute of Technology, Navi Mumbai

Aug 2019 – Jun 2022